

SQL: Advanced Queries, Assertions, Triggers, and Views

NULLS IN SQL QUERIES

- SQL allows queries that check if a value is **NULL** (missing or undefined or not applicable)
- SQL uses **IS** or **IS NOT** to compare NULLs because it considers each NULL value distinct from other NULL values, so *equality comparison is not appropriate*.
- Query 14: Retrieve the names of all employees who do not have supervisors.

Q14: SELECT FNAME, LNAME
 FROM EMPLOYEE
 WHERE SUPERSSN IS NULL

- Note: If a join condition is specified, tuples with NULL values for the join attributes are not included in the result

NESTING OF QUERIES

- A complete SELECT query, called a *nested query*, can be specified within the WHERE-clause of another query, called the *outer query*
 - Many of the previous queries can be specified in an alternative form using nesting
- Query 1: Retrieve the name and address of all employees who work for the 'Research' department.

Q1:SELECT	FNAME, LNAME, ADDRESS
FROM	EMPLOYEE
WHERE	DNO IN (SELECT DNUMBER
FROM	DEPARTMENT
WHERE	DNAME='Research')

NESTING OF QUERIES (contd.)

- The nested query selects the number of the 'Research' department
- The outer query select an EMPLOYEE tuple if its DNO value is in the result of either nested query
- The comparison operator **IN** compares a value v with a set (or multi-set) of values V , and evaluates to TRUE if v is one of the elements in V
- In general, we can have several levels of nested queries
- A reference to an *unqualified attribute* refers to the relation declared in the *innermost nested query*
- In this example, the nested query is *not correlated* with the outer query

CORRELATED NESTED QUERIES

- If a condition in the WHERE-clause of a *nested query* references an attribute of a relation declared in the *outer query*, the two queries are said to be *correlated*
 - The result of a correlated nested query is different for each tuple (or combination of tuples) of the relation(s) the outer query
- Query 12: Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
Q12: SELECT      E.FNAME, E.LNAME
      FROM        EMPLOYEE AS E
      WHERE       E.SSN IN
                  (SELECT      ESSN
                   FROM        DEPENDENT
                   WHERE       ESSN=E.SSN AND
                              E.FNAME=DEPENDENT_NAME)
```

CORRELATED NESTED QUERIES (contd.)

- In Q12, the nested query has a different result in the outer query
- A query written with nested SELECT... FROM... WHERE... blocks and using the = or IN comparison operators can ***always*** be expressed as a single block query. For example, Q12 may be written as in Q12A

```
Q12A:  SELECT      E.FNAME, E.LNAME
        FROM        EMPLOYEE E, DEPENDENT D
        WHERE       E.SSN=D.ESSN AND
                   E.FNAME=D.DEPENDENT_NAME
```

CORRELATED NESTED QUERIES (contd.)

- The original SQL as specified for SYSTEM R also had a **CONTAINS** comparison operator, which is used in conjunction with nested correlated queries
 - This operator was *dropped from the language*, possibly because of the difficulty in implementing it efficiently
 - Most implementations of SQL do not have this operator
 - The CONTAINS operator compares *two sets of values*, and returns TRUE if one set contains all values in the other set
 - Reminiscent of the division operation of algebra

CORRELATED NESTED QUERIES

(contd.)

- Query 3: Retrieve the name of each employee who works on all the projects controlled by department number 5.

```
Q3:      SELECT      FNAME, LNAME
          FROM        EMPLOYEE
          WHERE (      (SELECT      PNO
                        FROM        WORKS_ON
                        WHERE        SSN=ESSN)
                    CONTAINS
                        (SELECT      PNUMBER
                        FROM        PROJECT
                        WHERE        DNUM=5) )
```


CORRELATED NESTED QUERIES (contd.)

- In Q3, the second nested query, which is *not correlated* with the outer query, retrieves the project numbers of all projects controlled by department 5
- The first nested query, which is correlated, retrieves the project numbers on which the employee works, which is *different for each employee tuple* because of the correlation

THE EXISTS FUNCTION

- EXISTS is used to check whether the result of a correlated nested query is empty (contains no tuples) or not
 - We can formulate Query 12 in an alternative form that uses EXISTS as Q12B

THE EXISTS FUNCTION (contd.)

- Query 12: Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
Q12B:      SELECT      FNAME, LNAME
            FROM        EMPLOYEE
            WHERE        EXISTS (SELECT      *
                                FROM        DEPENDENT
                                WHERE        SSN=ESSN
                                AND
                                FNAME=DEPENDENT_NAME)
```

THE EXISTS FUNCTION (contd.)

- Query 6: Retrieve the names of employees who have no dependents.

```
Q6:      SELECT      FNAME, LNAME
          FROM        EMPLOYEE
          WHERE        NOT EXISTS (SELECT      *
                                   FROM        DEPENDENT
                                   WHERE        SSN=ESSN)
```

- In Q6, the correlated nested query retrieves all DEPENDENT tuples related to an EMPLOYEE tuple. If *none exist*, the EMPLOYEE tuple is selected
 - EXISTS is necessary for the expressive power of SQL

EXPLICIT SETS

- It is also possible to use an **explicit (enumerated) set of values** in the WHERE-clause rather than a nested query
- Query 13: Retrieve the social security numbers of all employees who work on project number 1, 2, or 3.

```
Q13:      SELECT      DISTINCT ESSN
           FROM        WORKS_ON
           WHERE        PNO IN (1, 2, 3)
```

Joined Relations Feature in SQL2

- Can specify a "joined relation" in the FROM-clause
 - Looks like any other relation but is the result of a join
 - Allows the user to specify different types of joins (regular "theta" JOIN, NATURAL JOIN, LEFT OUTER JOIN, RIGHT OUTER JOIN, CROSS JOIN, etc)

Joined Relations Feature in SQL2 (contd.)

- Examples:

Q8:SELECT	E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM	EMPLOYEE E S
WHERE	E.SUPERSSN=S.SSN

- can be written as:

Q8:SELECT	E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM	(EMPLOYEE E LEFT OUTER JOIN
	EMPLOYEE S ON E.SUPERSSN=S.SSN)

Joined Relations Feature in SQL2 (contd.)

- Examples:

```
Q1:SELECT      FNAME, LNAME, ADDRESS
      FROM EMPLOYEE, DEPARTMENT
      WHERE      DNAME='Research' AND DNUMBER=DNO
```

- could be written as:

```
Q1:SELECT      FNAME, LNAME, ADDRESS
      FROM      (EMPLOYEE JOIN DEPARTMENT
                  ON DNUMBER=DNO)
      WHERE      DNAME='Research'
```

- or as:

```
Q1:SELECT      FNAME, LNAME, ADDRESS
      FROM      (EMPLOYEE NATURAL JOIN
                  DEPARTMENT
                  AS DEPT(DNAME, DNO, MSSN, MSDATE))
      WHERE      DNAME='Research'
```


Joined Relations Feature in SQL2 (contd.)

- Another Example: Q2 could be written as follows; this illustrates multiple joins in the joined tables

```
Q2:      SELECT      PNUMBER, DNUM, LNAME,  
                      BDATE, ADDRESS  
FROM      (PROJECT JOIN  
           DEPARTMENT ON  
           DNUM=DNUMBER) JOIN  
           EMPLOYEE ON  
           MGRSSN=SSN) )  
WHERE     PLOCATION='Stafford'
```

AGGREGATE FUNCTIONS

- Include **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**
- Query 15: Find the maximum salary, the minimum salary, and the average salary among all employees.

```
Q15:      SELECT      MAX(SALARY),  
                MIN(SALARY), AVG(SALARY)  
          FROM      EMPLOYEE
```

- Some SQL implementations *may not allow more than one function* in the SELECT-clause

AGGREGATE FUNCTIONS (contd.)

- Query 16: Find the maximum salary, the minimum salary, and the average salary among employees who work for the 'Research' department.

```
Q16:      SELECT      MAX(SALARY),  
                  MIN(SALARY), AVG(SALARY)  
          FROM      EMPLOYEE, DEPARTMENT  
          WHERE      DNO=DNUMBER AND  
                  DNAME='Research'
```

AGGREGATE FUNCTIONS (contd.)

- Queries 17 and 18: Retrieve the total number of employees in the company (Q17), and the number of employees in the 'Research' department (Q18).

Q17: SELECT COUNT (*)
 FROM EMPLOYEE

Q18: SELECT COUNT (*)
 FROM EMPLOYEE, DEPARTMENT
 WHERE DNO=DNUMBER AND
 DNAME='Research'

GROUPING

- In many cases, we want to apply the aggregate functions to *subgroups of tuples* in a relation
- Each subgroup of tuples consists of the set of tuples that have the *same value* for the *grouping attribute(s)*
- The function is applied to each subgroup independently
- SQL has a **GROUP BY**-clause for specifying the grouping attributes, which *must also appear in the SELECT-clause*

GROUPING (contd.)

- Query 20: For each department, retrieve the department number, the number of employees in the department, and their average salary.

```
Q20:      SELECT      DNO, COUNT (*), AVG (SALARY)
           FROM        EMPLOYEE
           GROUP BY    DNO
```

- In Q20, the EMPLOYEE tuples are divided into groups-
 - Each group having the same value for the grouping attribute DNO
- The COUNT and AVG functions are applied to each such group of tuples separately
- The SELECT-clause includes only the grouping attribute and the functions to be applied on each group of tuples
- A join condition can be used in conjunction with grouping

GROUPING (contd.)

- Query 21: For each project, retrieve the project number, project name, and the number of employees who work on that project.

```
Q21:      SELECT      PNUMBER, PNAME, COUNT (*)
           FROM        PROJECT, WORKS_ON
           WHERE       PNUMBER=PNO
           GROUP BY    PNUMBER, PNAME
```

- In this case, the grouping and functions are applied after the joining of the two relations

THE HAVING-CLAUSE

- Sometimes we want to retrieve the values of these functions for only those *groups that satisfy certain conditions*
- The **HAVING**-clause is used for specifying a selection condition on groups (rather than on individual tuples)

THE HAVING-CLAUSE (contd.)

- Query 22: For each project *on which more than two employees work*, retrieve the project number, project name, and the number of employees who work on that project.

```
Q22:      SELECT      PNUMBER, PNAME,
                  COUNT(*)
            FROM        PROJECT, WORKS_ON
            WHERE        PNUMBER=PNO
            GROUP BY    PNUMBER, PNAME
            HAVING      COUNT (*) > 2
```

Summary of SQL Queries

- A query in SQL can consist of up to six clauses, but only the first two, **SELECT** and **FROM**, are mandatory. The clauses are specified in the following order:

SELECT	<attribute list>
FROM	<table list>
[WHERE	<condition>]
[GROUP BY	<grouping attribute(s)>]
[HAVING	<group condition>]
[ORDER BY	<attribute list>]

Summary of SQL Queries (contd.)

- The SELECT-clause lists the attributes or functions to be retrieved
- The FROM-clause specifies all relations (or aliases) needed in the query but not those needed in nested queries
- The WHERE-clause specifies the conditions for selection and join of tuples from the relations specified in the FROM-clause
- GROUP BY specifies grouping attributes
- HAVING specifies a condition for selection of groups
- ORDER BY specifies an order for displaying the result of a query
 - A query is evaluated by first applying the WHERE-clause, then GROUP BY and HAVING, and finally the SELECT-clause

Constraints as Assertions

- General constraints: constraints that do not fit in the basic SQL categories (presented in chapter 8)
- Mechanism: **CREATE ASSERTION**
 - Components include:
 - a constraint name,
 - followed by `CHECK`,
 - followed by a condition

Assertions: An Example

- “The salary of an employee must not be greater than the salary of the manager of the department that the employee works for”

CREAT ASSERTION **SALARY_CONSTRAINT**

CHECK

```
(NOT EXISTS (SELECT *  
              FROM EMPLOYEE E, EMPLOYEE M,  
                   DEPARTMENT D  
              WHERE E.SALARY > M.SALARY AND  
                    E.DNO=D.NUMBER AND  
                    D.MGRSSN=M.SSN) )
```

constraint
name,
CHECK,
condition

Using General Assertions

- Specify a query that violates the condition; include inside a `NOT EXISTS` clause
- Query result must be empty
 - if the query result is not empty, the assertion has been violated

SQL Triggers

- Objective: to monitor a database and take initiate action when a condition occurs
- Triggers are expressed in a syntax similar to assertions and include the following:
 - Event
 - Such as an insert, deleted, or update operation
 - Condition
 - Action
 - To be taken when the condition is satisfied

SQL Triggers: An Example

- A trigger to compare an employee's salary to his/her supervisor during insert or update operations:

```
CREATE TRIGGER SALARY_VIOLATION
BEFORE INSERT OR UPDATE OF
    SALARY, SUPERVISOR_SSN ON EMPLOYEE
FOR EACH ROW
    WHEN
        (NEW.SALARY > (SELECT SALARY FROM EMPLOYEE
                        WHERE SSN=NEW.SUPERVISOR_SSN) )
    INFORM_SUPERVISOR (NEW.SUPERVISOR_SSN, NEW.SSN) ;
```


Views in SQL

- A view is a “virtual” table that is derived from other tables
- Allows for limited update operations
 - Since the table may not physically be stored
- Allows full query operations
- A convenience for expressing certain operations

Specification of Views

- SQL command: **CREATE VIEW**
 - a table (view) name
 - a possible list of attribute names (for example, when arithmetic operations are specified or when we want the names to be different from the attributes in the base relations)
 - a query to specify the table contents

SQL Views: An Example

- Specify a different WORKS_ON table

```
CREATE VIEW WORKS_ON_NEW AS
SELECT FNAME, LNAME, PNAME, HOURS
FROM EMPLOYEE, PROJECT, WORKS_ON
WHERE SSN=ESSN AND PNO=PNUMBER
GROUP BY PNAME;
```

Using a Virtual Table

- We can specify SQL queries on a newly create table (view):

```
SELECT FNAME, LNAME  
      FROM WORKS_ON_NEW  
      WHERE PNAME='Seena' ;
```

- When no longer needed, a view can be dropped:

```
DROP WORKS_ON_NEW;
```

Efficient View Implementation

- Query modification:
 - Present the view query in terms of a query on the underlying base tables
- Disadvantage:
 - Inefficient for views defined via complex queries
 - Especially if additional queries are to be applied to the view within a short time period

Efficient View Implementation

- View materialization:
 - Involves physically creating and keeping a temporary table
- Assumption:
 - Other queries on the view will follow
- Concerns:
 - Maintaining correspondence between the base table and the view when the base table is updated
- Strategy:
 - Incremental update

Update Views

- Update on a single view without aggregate operations:
 - Update may map to an update on the underlying base table
- Views involving joins:
 - An update *may* map to an update on the underlying base relations
 - Not always possible

Un-updatable Views

- Views defined using groups and aggregate functions are not updateable
- Views defined on multiple tables using joins are generally not updateable
- **WITH CHECK OPTION**: must be added to the definition of a view if the view is to be updated
 - To allow check for updatability and to plan for an execution strategy